

# Yelpers in The Great Recession

2/7/2017

## outline

intro data set explanations yelp data fred analysis yelp users time series yelp reviews time series yelp review stars time series fred data for recession analysis yelp stock data conclusion

## Introduction

Hypothesis: yelpers

## Data Sets

yelp fred yahoo finance bea

## Motivation

If restaurant Yelpers' behavior during and around The Great Recession period can be modeled and connected with restaurant expenditures, then there can be conclusions about Yelpers' behavior. If recessionary behavior can be understood, then restaurants can react accordingly.

```
# SETUP
setwd("C:/cygwin64/home/Lester/yelp_challenge_9")

# load data example json_file =
# file('yelp_academic_dataset_checkin.json') json_data =
# jsonlite::stream_in(json_file) head(json_data)
# length(json_data$business_id)

load_json = function(filename) {
  json_file = file(filename)
  json_data = jsonlite::stream_in(json_file)
  return(json_data)
}

# business = load_json('yelp_academic_dataset_business.json')
# review = load_json('yelp_academic_dataset_review.json')
# checkin = load_json('yelp_academic_dataset_checkin.json')
# tip = load_json('yelp_academic_dataset_tip.json') user =
# load_json('yelp_academic_dataset_user.json')

remove_lists_from_df = function(df) {
  i = 1
  while (i <= length(df)) {
    if (class(df[, i]) == "list") {
      df[i] = sapply(df[, i], paste, collapse = "|")
    }
  }
}
```

```

    }
    i = i + 1
  }

  return(df)
}

add_recession_dummy = function(l) {
  rec = c()
  for (i in 1:length(l)) {


    if (l[i] >= as.Date("2007-12-01") & l[i] <= as.Date("2009-07-01")) {
      rec = c(rec, 1)
    } else {
      rec = c(rec, 0)
    }
  }

  return(rec)
}

```

## user data

```
plot(stock_users, type = "l")
```




```
yelp_stock = get.hist.quote("YELP", end = "2017-01-30", quote = "AdjClose",  
    compression = "m")
```

```
## time series starts 2012-03-02  
## time series ends 2017-01-03
```

```
# yelp_stock
```

```
plot(yelp_stock)
```



```
test_stationary = function(t) {
  print(kpss.test(t))
  print(adf.test(t))
}

test_stationary(ts(stock_users$coredata.ts_m., start = c(2012,
3), freq = 12))

## Warning in kpss.test(t): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 1.1995, Truncation lag parameter = 1, p-value = 0.01
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -1.0302, Lag order = 3, p-value = 0.9257
## alternative hypothesis: stationary
test_stationary(yelp_stock)

##
## KPSS Test for Level Stationarity
##
```

```

## data:  t
## KPSS Level = 0.53728, Truncation lag parameter = 1, p-value =
## 0.03327
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -1.4945, Lag order = 3, p-value = 0.7789
## alternative hypothesis: stationary

log_user_growth = as.data.frame(diff(log(stock_users$coredata.ts_m.)))

log_yelp_growth = as.data.frame(diff(log(yelp_stock)))

# log_user_growth log_yelp_growth

ts_users = ts(log_user_growth, start = c(2012, 4), freq = 12)
test_stationary(ts_users)


## Warning in kpss.test(t): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 0.8162, Truncation lag parameter = 1, p-value = 0.01
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -2.5657, Lag order = 3, p-value = 0.3463
## alternative hypothesis: stationary

ts_yelp = ts(log_yelp_growth, start = c(2012, 4), freq = 12)
test_stationary(ts_yelp)


## Warning in kpss.test(t): p-value greater than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 0.13073, Truncation lag parameter = 1, p-value = 0.1
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -3.9873, Lag order = 3, p-value = 0.01641
## alternative hypothesis: stationary

```

```
plot(ts_users, col = "blue")
lines(ts_yelp, col = "red")
```



```
combined = cbind(ts_yelp, ts_users)
select = VARselect(combined, lag.max = 12, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)
vm = VAR(combined, p = select$select[1])
plot(vm$y)
```



```
summary(vm)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: ts_yelp, ts_users
## Deterministic variables: const
## Sample size: 46
## Log Likelihood: 104.197
## Roots of the characteristic polynomial:
## 1.121 0.9968 0.9968 0.9859 0.9859 0.9739 0.9739 0.9716 0.9716 0.9556 0.9556 0.9471 0.9423 0.9423 0.9
## Call:
## VAR(y = combined, p = select$select[1])
##
##
## Estimation results for equation ts_yelp:
## =====
## ts_yelp = ts_yelp.l1 + ts_users.l1 + ts_yelp.l2 + ts_users.l2 + ts_yelp.l3 + ts_users.l3 + ts_yelp.l
##
##           Estimate Std. Error t value Pr(>|t|)
## ts_yelp.l1    0.22443    0.24350   0.922   0.367
## ts_users.l1    0.27950    0.43228   0.647   0.525
## ts_yelp.l2   -0.02723    0.26692  -0.102   0.920
## ts_users.l2   -0.12729    0.44781  -0.284   0.779
## ts_yelp.l3    0.06446    0.25133   0.256   0.800
## ts_users.l3   -0.11399    0.47105  -0.242   0.811
```

```

## ts_yelp.l4    -0.15175    0.26027   -0.583    0.566
## ts_users.l4    0.41380    0.39727    1.042    0.309
## ts_yelp.l5     0.34786    0.24875    1.398    0.177
## ts_users.l5   -0.04790    0.41595   -0.115    0.909
## ts_yelp.l6     0.23889    0.29288    0.816    0.424
## ts_users.l6   -0.01589    0.38874   -0.041    0.968
## ts_yelp.l7    -0.07162    0.29937   -0.239    0.813
## ts_users.l7   -0.33711    0.38189   -0.883    0.387
## ts_yelp.l8     0.01820    0.29149    0.062    0.951
## ts_users.l8   -0.22491    0.40855   -0.551    0.588
## ts_yelp.l9     0.07345    0.27828    0.264    0.794
## ts_users.l9    0.40564    0.37089    1.094    0.286
## ts_yelp.l10   -0.15597    0.24322   -0.641    0.528
## ts_users.l10   0.31981    0.38672    0.827    0.418
## ts_yelp.l11   -0.16642    0.22890   -0.727    0.475
## ts_users.l11  -0.26574    0.41547   -0.640    0.529
## ts_yelp.l12   -0.13732    0.23143   -0.593    0.559
## ts_users.l12   0.05773    0.44119    0.131    0.897
## const         0.01366    0.03194    0.428    0.673
##
##
## Residual standard error: 0.1904 on 21 degrees of freedom
## Multiple R-Squared: 0.411,   Adjusted R-squared: -0.2621
## F-statistic: 0.6106 on 24 and 21 DF,  p-value: 0.8779
##
##
## Estimation results for equation ts_users:
## =====
## ts_users = ts_yelp.l1 + ts_users.l1 + ts_yelp.l2 + ts_users.l2 + ts_yelp.l3 + ts_users.l3 + ts_yelp.l4 + ts_users.l4 + ts_yelp.l5 + ts_users.l5 + ts_yelp.l6 + ts_users.l6 + ts_yelp.l7 + ts_users.l7 + ts_yelp.l8 + ts_users.l8 + ts_yelp.l9 + ts_users.l9 + ts_yelp.l10 + ts_users.l10 + ts_yelp.l11 + ts_users.l11 + ts_yelp.l12 + ts_users.l12
##
##
##           Estimate Std. Error t value Pr(>|t|)
## ts_yelp.l1   -0.158567    0.103539   -1.531  0.14058
## ts_users.l1    0.062517    0.183814    0.340  0.73715
## ts_yelp.l2   -0.034134    0.113501   -0.301  0.76657
## ts_users.l2    0.353135    0.190419    1.855  0.07776 .
## ts_yelp.l3    0.004447    0.106872    0.042  0.96720
## ts_users.l3    0.185054    0.200299    0.924  0.36604
## ts_yelp.l4   -0.072499    0.110674   -0.655  0.51954
## ts_users.l4   -0.073652    0.168926   -0.436  0.66728
## ts_yelp.l5   -0.163585    0.105773   -1.547  0.13691
## ts_users.l5    0.041078    0.176871    0.232  0.81859
## ts_yelp.l6    0.235357    0.124538    1.890  0.07266 .
## ts_users.l6    0.140393    0.165300    0.849  0.40528
## ts_yelp.l7   -0.105272    0.127298   -0.827  0.41755
## ts_users.l7    0.428628    0.162388    2.640  0.01533 *
## ts_yelp.l8    0.053978    0.123949    0.435  0.66765
## ts_users.l8    0.031861    0.173722    0.183  0.85624
## ts_yelp.l9   -0.015866    0.118330   -0.134  0.89462
## ts_users.l9   -0.097292    0.157710   -0.617  0.54393
## ts_yelp.l10   0.053363    0.103420    0.516  0.61126
## ts_users.l10   0.264388    0.164440    1.608  0.12281
## ts_yelp.l11   0.092298    0.097331    0.948  0.35377
## ts_users.l11   0.439468    0.176665    2.488  0.02134 *
## ts_yelp.l12  -0.029449    0.098410   -0.299  0.76769

```



```
## ts_users.l12  0.662858  0.187602  3.533  0.00197 **
## const        -0.009345  0.013580 -0.688  0.49886
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.08095 on 21 degrees of freedom
## Multiple R-Squared:  0.798,    Adjusted R-squared:  0.5672
## F-statistic: 3.457 on 24 and 21 DF,  p-value: 0.002768
##
##
## Covariance matrix of residuals:
##          ts_yelp ts_users
## ts_yelp  0.036241 0.007759
## ts_users 0.007759 0.006553
##
## Correlation matrix of residuals:
##          ts_yelp ts_users
## ts_yelp  1.0000  0.5035
## ts_users  0.5035  1.0000

grangertest(ts_users ~ ts_yelp)

## Granger causality test
##
## Model 1: ts_users ~ Lags(ts_users, 1:1) + Lags(ts_yelp, 1:1)
## Model 2: ts_users ~ Lags(ts_users, 1:1)
##   Res.Df Df       F Pr(>F)
## 1      54
## 2      55 -1 0.2396 0.6265

sp500 = get.hist.quote("^GSPC", quote = "AdjClose", compression = "m",
  start = as.Date("2012-03-01", "%Y-%m-%d"), end = as.Date("2017-01-30",
    "%Y-%m-%d"))

## time series ends  2017-01-03

# sp500
test_stationary(sp500)

## Warning in kpss.test(t): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 2.7456, Truncation lag parameter = 1, p-value = 0.01
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -1.3383, Lag order = 3, p-value = 0.842
## alternative hypothesis: stationary
```


```

log_sp500_growth = as.data.frame(diff(log(sp500)))
ts_sp500 = ts(log_sp500_growth, start = c(2012, 4), freq = 12)
# ts_sp500
test_stationary(ts_sp500)

## Warning in kpss.test(t): p-value greater than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 0.096794, Truncation lag parameter = 1, p-value = 0.1
## Warning in adf.test(t): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -5.2647, Lag order = 3, p-value = 0.01
## alternative hypothesis: stationary
# ts_yelp ts_sp500

plot(ts_users, col = "blue")
lines(ts_yelp, col = "red")

```




```
combined = cbind(ts_yelp, ts_sp500)

select = VARselect(combined, lag.max = 12, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)

vm = VAR(combined, p = select$select[1])

plot(vm$y)
```



```
summary(vm)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: ts_yelp, ts_sp500
## Deterministic variables: const
## Sample size: 57
## Log Likelihood: 146.967
## Roots of the characteristic polynomial:
## 0.261 0.09083
## Call:
## VAR(y = combined, p = select$select[1])
##
##
## Estimation results for equation ts_yelp:
## =====
```

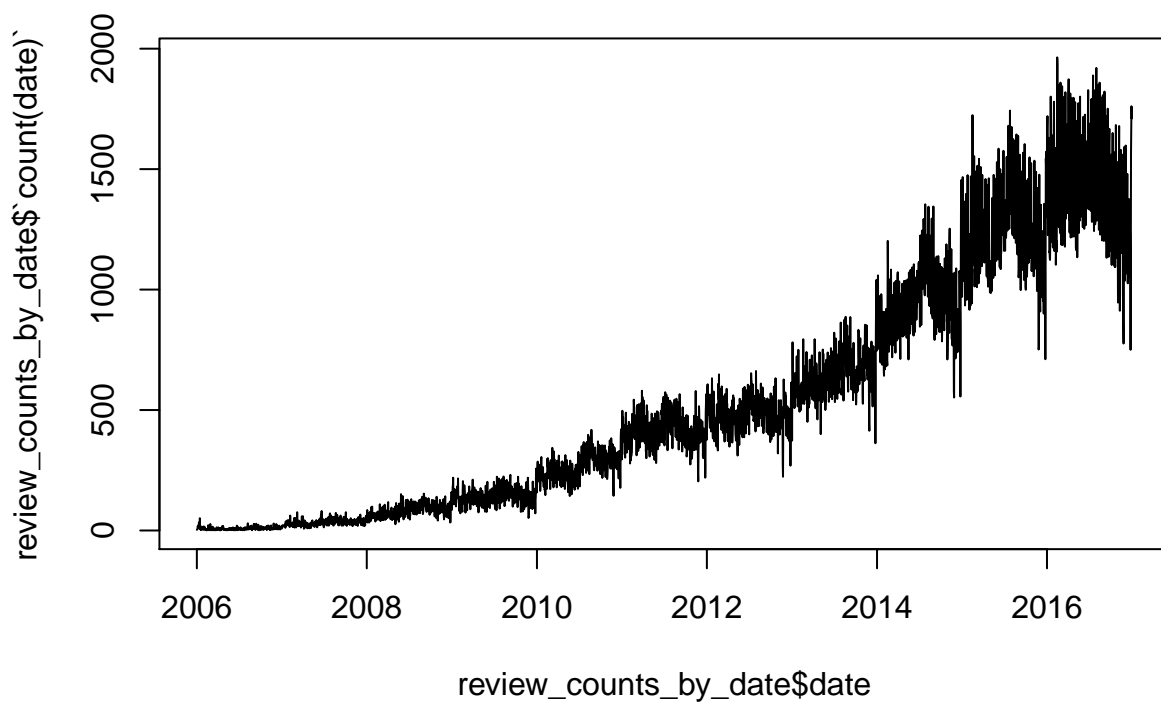
```
## ts_yelp = ts_yelp.l1 + ts_sp500.l1 + const
##
##           Estimate Std. Error t value Pr(>|t|)
## ts_yelp.l1 0.062844 0.144375 0.435 0.665
## ts_sp500.l1 0.524613 0.838397 0.626 0.534
## const      0.005891 0.023856 0.247 0.806
##
##
## Residual standard error: 0.1732 on 54 degrees of freedom
## Multiple R-Squared: 0.01655, Adjusted R-squared: -0.01987
## F-statistic: 0.4545 on 2 and 54 DF, p-value: 0.6372
##
##
## Estimation results for equation ts_sp500:
## =====
## ts_sp500 = ts_yelp.l1 + ts_sp500.l1 + const
##
##           Estimate Std. Error t value Pr(>|t|)
## ts_yelp.l1 0.017279 0.024603 0.702 0.4855
## ts_sp500.l1 -0.233048 0.142871 -1.631 0.1087
## const      0.010378 0.004065 2.553 0.0135 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.02952 on 54 degrees of freedom
## Multiple R-Squared: 0.04714, Adjusted R-squared: 0.01185
## F-statistic: 1.336 on 2 and 54 DF, p-value: 0.2715
##
##
## Covariance matrix of residuals:
##      ts_yelp ts_sp500
## ts_yelp 0.030002 0.0020338
## ts_sp500 0.002034 0.0008712
##
## Correlation matrix of residuals:
##      ts_yelp ts_sp500
## ts_yelp 1.0000 0.3978
## ts_sp500 0.3978 1.0000
```

Growth rate of users do not have an effect on the stock value of Yelp - no correlation to stock market. good!  
This eliminates the possibility that effects are due to yelp as a company doing good or bad.

should i examine review scores + users + etc and its relation to s and p500 during recession? or just move on to gdp and recession

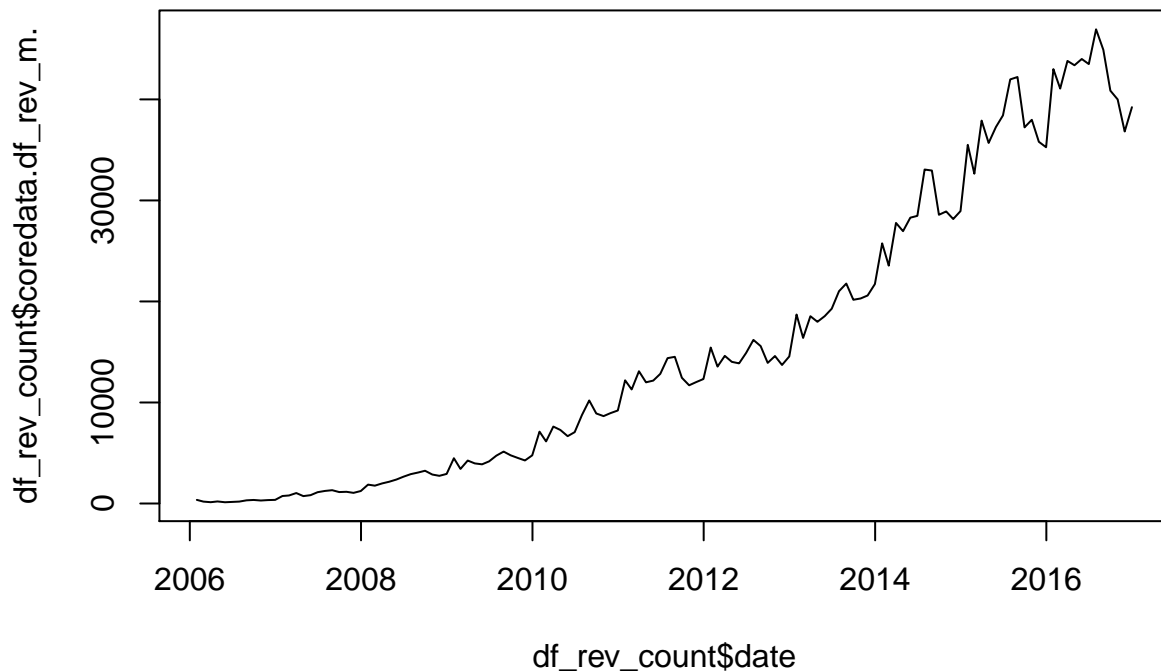
## review scores and gdp

```
plot(review_counts_by_date$date, review_counts_by_date$count(date),
     type = "l")
```



```
# convert to monthly
reviews_by_date = xts(review_counts_by_date$count(date), as.Date(review_counts_by_date$date,
"%Y-%m-%d"))
df_rev_m = apply.monthly(reviews_by_date, sum)
df_rev_count = data.frame(date = index(df_rev_m), coredata(df_rev_m))
# df_rev_count

plot(df_rev_count$date, df_rev_count$coredata.df_rev_m., type = "l")
```




```
# review growth rates/new user growth rates
log_rev_count = diff(log(df_rev_count$coredata.df_rev_m.))
# log_rev_count[1]=NA
log_rev_count = na.omit(log_rev_count)
log_rev_count = ts(log_rev_count, start = c(2006, 2), freq = 12)

# ts_m
log_user_count = diff(log(ts_m[, 1]))
log_user_count = na.omit(log_user_count)
log_user_count = ts(log_user_count, start = c(2006, 2), freq = 12)


plot(log_rev_count, type = "l", main = "growth rate of user reviews and accounts")
lines(log_user_count[, 1], col = "red")
```

## growth rate of user reviews and accounts



```
# do a var model between growth rate of users revs and  
# accounts  
  
# create var of growth rates  
rates_combined = cbind(log_rev_count, log_user_count)  
select = VARselect(rates_combined, lag.max = 12, type = c("const",  
  "trend", "both", "none"), season = NULL, exogen = NULL)  
vm_rates = VAR(rates_combined, select$select[1])  
plot(vm_rates$y)
```

## vm\_rates\$y



```
summary(vm_rates)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: log_rev_count, log_user_count
## Deterministic variables: const
## Sample size: 119
## Log Likelihood: 297.714
## Roots of the characteristic polynomial:
## 1.048 0.9737 0.9737 0.9716 0.9716 0.97 0.97 0.9687 0.9581 0.9581 0.9101 0.9101 0.8859 0.8859 0.857
## Call:
## VAR(y = rates_combined, p = select$select[1])
##
##
## Estimation results for equation log_rev_count:
## =====
## log_rev_count = log_rev_count.l1 + log_user_count.l1 + log_rev_count.l2 + log_user_count.l2 + log_rev_count.l3 + log_user_count.l3
##
##              Estimate Std. Error t value Pr(>|t|)
## log_rev_count.l1 -0.447591 0.102482 -4.367 3.23e-05 ***
## log_user_count.l1 0.073934 0.092299 0.801 0.425138
## log_rev_count.l2 -0.140889 0.109751 -1.284 0.202398
## log_user_count.l2 0.208054 0.097639 2.131 0.035711 *
## log_rev_count.l3 0.024924 0.105031 0.237 0.812936
## log_user_count.l3 -0.053311 0.095380 -0.559 0.577540
```



```

## log_rev_count.l4      0.210029    0.098349    2.136 0.035316 *
## log_user_count.l4    -0.405138    0.096096   -4.216 5.72e-05 ***
## log_rev_count.l5      0.005835    0.095989    0.061 0.951655
## log_user_count.l5    -0.071718    0.102943   -0.697 0.487721
## log_rev_count.l6      0.113966    0.094900    1.201 0.232807
## log_user_count.l6     0.108670    0.101916    1.066 0.289033
## log_rev_count.l7      0.115462    0.094805    1.218 0.226318
## log_user_count.l7     0.124041    0.103038    1.204 0.231678
## log_rev_count.l8     -0.084966    0.095834   -0.887 0.377561
## log_user_count.l8     0.003598    0.103881    0.035 0.972444
## log_rev_count.l9     -0.243305    0.090552   -2.687 0.008528 **
## log_user_count.l9     0.122566    0.101162    1.212 0.228713
## log_rev_count.l10    -0.184751    0.095714   -1.930 0.056592 .
## log_user_count.l10    0.299530    0.100755    2.973 0.003748 **
## log_rev_count.l11    -0.132124    0.086128   -1.534 0.128376
## log_user_count.l11    0.214219    0.101571    2.109 0.037598 *
## log_rev_count.l12     0.092019    0.066103    1.392 0.167189
## log_user_count.l12    0.373805    0.096160    3.887 0.000189 ***
## const                 0.032366    0.013296    2.434 0.016811 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.07862 on 94 degrees of freedom
## Multiple R-Squared:  0.6877, Adjusted R-squared:  0.6079
## F-statistic: 8.623 on 24 and 94 DF, p-value: 7.305e-15
##
##
## Estimation results for equation log_user_count:
## =====
## log_user_count = log_rev_count.l11 + log_user_count.l11 + log_rev_count.l12 + log_user_count.l12 + log_r
##
##
##              Estimate Std. Error t value Pr(>|t|)
## log_rev_count.l11 -0.228340    0.108927  -2.096  0.03874 *
## log_user_count.l11  0.085515    0.098103   0.872  0.38560
## log_rev_count.l12  0.109457    0.116653   0.938  0.35049
## log_user_count.l12  0.032154    0.103779   0.310  0.75738
## log_rev_count.l13  0.155007    0.111635   1.389  0.16826
## log_user_count.l13  0.086143    0.101378   0.850  0.39764
## log_rev_count.l14  0.212001    0.104533   2.028  0.04538 *
## log_user_count.l14 -0.139470    0.102139  -1.365  0.17536
## log_rev_count.l15 -0.003095    0.102025  -0.030  0.97586
## log_user_count.l15  0.064438    0.109416   0.589  0.55732
## log_rev_count.l16 -0.001041    0.100868  -0.010  0.99179
## log_user_count.l16  0.187274    0.108324   1.729  0.08712 .
## log_rev_count.l17 -0.085882    0.100767  -0.852  0.39622
## log_user_count.l17  0.260097    0.109517   2.375  0.01958 *
## log_rev_count.l18 -0.212682    0.101861  -2.088  0.03951 *
## log_user_count.l18  0.012833    0.110413   0.116  0.90772
## log_rev_count.l19 -0.310263    0.096246  -3.224  0.00174 **
## log_user_count.l19  0.237115    0.107524   2.205  0.02988 *
## log_rev_count.l110 -0.282063    0.101732  -2.773  0.00671 **
## log_user_count.l110  0.349787    0.107091   3.266  0.00152 **
## log_rev_count.l111 -0.161645    0.091543  -1.766  0.08068 .

```

```

## log_user_count.l11 0.258147 0.107957 2.391 0.01879 *
## log_rev_count.l12 0.256980 0.070259 3.658 0.00042 ***
## log_user_count.l12 0.257381 0.102207 2.518 0.01348 *
## const -0.008062 0.014132 -0.570 0.56971
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.08356 on 94 degrees of freedom
## Multiple R-Squared: 0.618, Adjusted R-squared: 0.5204
## F-statistic: 6.336 on 24 and 94 DF, p-value: 3.029e-11
##
##
## Covariance matrix of residuals:
##          log_rev_count log_user_count
## log_rev_count 0.006181 0.002504
## log_user_count 0.002504 0.006982
##
## Correlation matrix of residuals:
##          log_rev_count log_user_count
## log_rev_count 1.0000 0.3811
## log_user_count 0.3811 1.0000

```

```

# do granger causality test
grangertest(log_rev_count ~ log_user_count, order = select$select[1])

```

```

## Granger causality test
##
## Model 1: log_rev_count ~ Lags(log_rev_count, 1:12) + Lags(log_user_count, 1:12)
## Model 2: log_rev_count ~ Lags(log_rev_count, 1:12)
##   Res.Df Df       F    Pr(>F)
## 1      94
## 2     106 -12 5.9487 1.318e-07 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

grangertest(log_user_count ~ log_rev_count, order = select$select[1])

```

```

## Granger causality test
##
## Model 1: log_user_count ~ Lags(log_user_count, 1:12) + Lags(log_rev_count, 1:12)
## Model 2: log_user_count ~ Lags(log_user_count, 1:12)
##   Res.Df Df       F    Pr(>F)
## 1      94
## 2     106 -12 3.4167 0.000354 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

user_rev_lm = lm(log_rev_count ~ log_user_count)
summary(user_rev_lm)

```

```

##
## Call:
## lm(formula = log_rev_count ~ log_user_count)
##

```


```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.68573 -0.04008  0.00958  0.04946  0.41635
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.02131    0.01183   1.801  0.0741 .
## log_user_count 0.77640    0.08763   8.860 5.44e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1341 on 129 degrees of freedom
## Multiple R-squared:  0.3783, Adjusted R-squared:  0.3735
## F-statistic: 78.51 on 1 and 129 DF,  p-value: 5.444e-15

# users and review#s granger cause each other, just use one

# stars by month
stars_by_date = xts(review_date_star$stars, as.Date(review_date_star$date,
"%Y-%m-%d"))
df_stars_m = apply.monthly(stars_by_date, sum)
df_stars = data.frame(date = index(df_stars_m), coredata(df_stars_m))

df_stars$avg = df_stars$coredata.df_stars_m./df_rev_count$coredata.df_rev_m.
# head(df_stars)

plot(df_stars$date, df_stars$avg, type = "l")
```



```
# evidence of recession in stars
stars_recession_dummy = add_recession_dummy(df_stars$date)
stars_avg = df_stars$avg
stars_lm = lm(stars_avg ~ stars_recession_dummy)
summary(stars_lm)

##
## Call:
## lm(formula = stars_avg ~ stars_recession_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.12291 -0.04541 -0.01606  0.04614  0.22330
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.712971   0.006261  593.052 < 2e-16 ***
## stars_recession_dummy -0.048566   0.016502  -2.943  0.00385 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06655 on 130 degrees of freedom
## Multiple R-squared:  0.06246,    Adjusted R-squared:  0.05525
## F-statistic: 8.661 on 1 and 130 DF,  p-value: 0.003851
```

```
test_stationary(df_stars$avg)
```

```
## Warning in kpss.test(t): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data:  t
## KPSS Level = 0.84648, Truncation lag parameter = 2, p-value = 0.01
##
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -1.9674, Lag order = 5, p-value = 0.5901
## alternative hypothesis: stationary
```

```
# convert to growth rates
```

```
df_stars_diff_log = as.data.frame(diff(log(df_stars$avg)))
df_stars_diff_log$date = df_stars$date[2:length(df_stars$date)]
```

```
df_stars_diff_log
```

```
##      diff(log(df_stars$avg))      date
## 1      -5.619865e-03 2006-02-28
## 2       1.923989e-02 2006-03-31
## 3       1.265143e-03 2006-04-29
## 4      -2.443172e-02 2006-05-31
## 5      -3.263090e-02 2006-06-30
## 6       1.963358e-02 2006-07-31
## 7      -1.612599e-02 2006-08-31
## 8       2.800839e-03 2006-09-30
## 9      -8.594358e-03 2006-10-31
## 10      1.400202e-02 2006-11-30
## 11     -1.633079e-02 2006-12-31
## 12      4.251363e-02 2007-01-31
## 13      3.453816e-04 2007-02-28
## 14     -1.845148e-02 2007-03-31
## 15     -1.075444e-02 2007-04-30
## 16      5.325714e-03 2007-05-31
## 17      5.534589e-03 2007-06-30
## 18     -1.201243e-02 2007-07-31
## 19      1.122151e-02 2007-08-31
## 20     -2.366109e-02 2007-09-30
## 21     -2.852157e-03 2007-10-31
## 22      3.710066e-02 2007-11-30
## 23     -3.173946e-02 2007-12-31
## 24     -2.795006e-03 2008-01-31
## 25      8.392467e-03 2008-02-29
## 26      8.993568e-04 2008-03-31
## 27     -7.813701e-03 2008-04-30
## 28     -3.015577e-03 2008-05-31
## 29     -8.163412e-03 2008-06-30
## 30     -3.424082e-03 2008-07-31
## 31     -6.027338e-03 2008-08-31
```

## 32	1.996352e-03	2008-09-30
## 33	-8.728518e-04	2008-10-31
## 34	-1.921703e-03	2008-11-30
## 35	5.202632e-03	2008-12-31
## 36	-1.341539e-03	2009-01-31
## 37	5.339377e-03	2009-02-28
## 38	-1.426190e-02	2009-03-31
## 39	1.348717e-02	2009-04-30
## 40	-7.656413e-03	2009-05-31
## 41	8.582641e-03	2009-06-30
## 42	7.958529e-03	2009-07-31
## 43	-1.406197e-02	2009-08-31
## 44	-6.107051e-03	2009-09-30
## 45	4.783419e-03	2009-10-31
## 46	-1.056028e-02	2009-11-30
## 47	-1.577807e-03	2009-12-31
## 48	7.166652e-03	2010-01-31
## 49	9.542976e-03	2010-02-28
## 50	-3.927418e-03	2010-03-31
## 51	4.125844e-03	2010-04-30
## 52	4.519455e-03	2010-05-31
## 53	-2.684434e-03	2010-06-30
## 54	1.074447e-02	2010-07-31
## 55	-5.534875e-03	2010-08-31
## 56	2.208555e-03	2010-09-30
## 57	-1.857518e-03	2010-10-31
## 58	3.900998e-04	2010-11-30
## 59	-2.770717e-03	2010-12-31
## 60	3.639439e-03	2011-01-31
## 61	-3.016042e-03	2011-02-28
## 62	-4.471740e-04	2011-03-31
## 63	1.281053e-04	2011-04-30
## 64	2.094403e-03	2011-05-31
## 65	4.849414e-03	2011-06-30
## 66	-1.482645e-03	2011-07-31
## 67	-1.330958e-02	2011-08-31
## 68	8.154915e-04	2011-09-30
## 69	4.060745e-03	2011-10-31
## 70	9.731713e-04	2011-11-30
## 71	-8.223230e-03	2011-12-31
## 72	6.484596e-03	2012-01-31
## 73	-2.251225e-03	2012-02-29
## 74	-5.706624e-03	2012-03-31
## 75	3.849061e-03	2012-04-30
## 76	3.306098e-03	2012-05-31
## 77	-2.592167e-03	2012-06-30
## 78	8.750600e-04	2012-07-31
## 79	2.467858e-03	2012-08-31
## 80	1.041843e-03	2012-09-30
## 81	-3.053989e-03	2012-10-31
## 82	1.161721e-03	2012-11-30
## 83	-6.995426e-03	2012-12-31
## 84	1.166825e-02	2013-01-31
## 85	-6.350921e-03	2013-02-28

```
## 86          1.749019e-04 2013-03-31
## 87        -7.492477e-03 2013-04-30
## 88          9.204614e-04 2013-05-31
## 89          1.312812e-02 2013-06-30
## 90          1.147936e-03 2013-07-31
## 91          4.490404e-03 2013-08-31
## 92        -6.214124e-03 2013-09-30
## 93          5.951240e-03 2013-10-31
## 94        -5.569303e-03 2013-11-30
## 95        -1.303337e-03 2013-12-31
## 96          1.348184e-02 2014-01-31
## 97          4.870529e-05 2014-02-28
## 98        -3.615021e-03 2014-03-31
## 99        -2.585917e-03 2014-04-30
## 100       -1.311054e-03 2014-05-31
## 101          6.378572e-03 2014-06-30
## 102          5.917766e-03 2014-07-31
## 103          5.419283e-04 2014-08-31
## 104       -1.002166e-02 2014-09-30
## 105          6.225047e-03 2014-10-31
## 106       -5.479281e-03 2014-11-30
## 107       -4.762909e-03 2014-12-31
## 108          1.292571e-02 2015-01-31
## 109       -7.020857e-03 2015-02-28
## 110       -2.731912e-03 2015-03-31
## 111          1.000075e-02 2015-04-30
## 112          2.110902e-03 2015-05-31
## 113          4.274085e-03 2015-06-30
## 114          7.234081e-03 2015-07-31
## 115       -3.011319e-03 2015-08-31
## 116       -1.848002e-03 2015-09-30
## 117          2.730000e-04 2015-10-31
## 118          1.595003e-03 2015-11-30
## 119       -8.085294e-03 2015-12-31
## 120          1.497098e-02 2016-01-31
## 121       -1.014415e-02 2016-02-29
## 122          2.609374e-03 2016-03-31
## 123          5.401924e-03 2016-04-30
## 124       -6.305894e-03 2016-05-31
## 125          8.415438e-03 2016-06-30
## 126       -1.068921e-03 2016-07-31
## 127       -2.013443e-03 2016-08-31
## 128          4.508334e-03 2016-09-30
## 129          1.314815e-04 2016-10-31
## 130       -4.937163e-03 2016-11-30
## 131       -7.075452e-03 2016-12-31
```

```
colnames(df_stars_diff_log) = c("avg", "date")

test_stationary(df_stars_diff_log$avg)
```

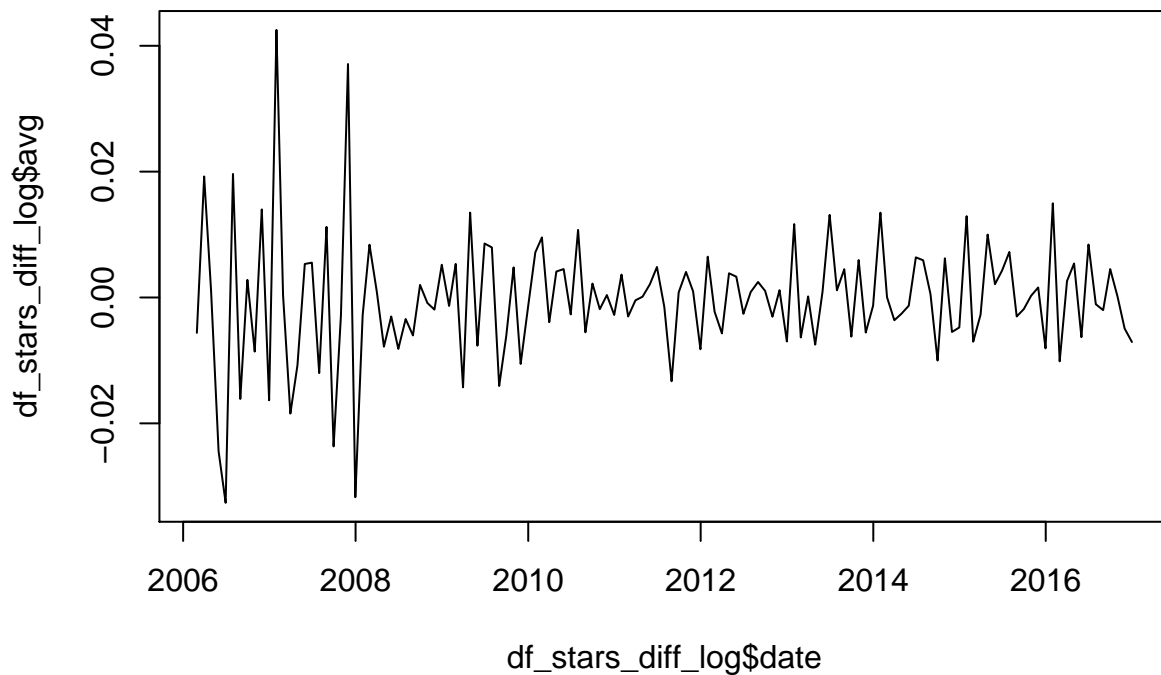
```
## Warning in kpss.test(t): p-value greater than printed p-value

##
## KPSS Test for Level Stationarity
```

```
##
## data:  t
## KPSS Level = 0.20563, Truncation lag parameter = 2, p-value = 0.1
## Warning in adf.test(t): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data:  t
## Dickey-Fuller = -5.5779, Lag order = 5, p-value = 0.01
## alternative hypothesis: stationary
stars_recession_diff_log_dummy = add_recession_dummy(df_stars_diff_log$date)

# stars_reg = lm(df_stars_diff_log$avg ~
# stars_recession_diff_log_dummy) summary(stars_reg)

plot(df_stars_diff_log$date, df_stars_diff_log$avg, type = "l")
```



```
stars_diff_log_avg = df_stars_diff_log$avg
stars_diff_log_lm = lm(stars_diff_log_avg ~ stars_recession_diff_log_dummy)
summary(stars_diff_log_lm)
```

```
##
## Call:
## lm(formula = stars_diff_log_avg ~ stars_recession_diff_log_dummy)
##
```



```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.032757 -0.005360  0.000147  0.004525  0.042388
##
## Coefficients:
##                  Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.0001261  0.0009446   0.133   0.894
## stars_recession_diff_log_dummy -0.0025015  0.0024803  -1.009   0.315
##
## Residual standard error: 0.009997 on 129 degrees of freedom
## Multiple R-squared:  0.007823,    Adjusted R-squared:  0.0001321
## F-statistic: 1.017 on 1 and 129 DF,  p-value: 0.3151
```


*# however, intuitively we should be looking at level, not  
# growth rates so lets detrend the data and season*

```
ts_stars = ts(df_stars$avg, start = c(2006, 2), freq = 12)
```


```
stars_tslm = tslm(ts_stars ~ trend + season)
summary(stars_tslm)
```

```
##
## Call:
## tslm(formula = ts_stars ~ trend + season)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.10883 -0.04940 -0.01198  0.04378  0.22099
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.677e+00  2.426e-02 151.563  <2e-16 ***
## trend       -2.068e-05  1.619e-04  -0.128   0.899
## season2      4.736e-02  3.015e-02   1.571   0.119
## season3      4.357e-02  3.014e-02   1.446   0.151
## season4      3.505e-02  3.013e-02   1.163   0.247
## season5      3.822e-02  3.012e-02   1.269   0.207
## season6      2.959e-02  3.011e-02   0.983   0.328
## season7      3.114e-02  3.011e-02   1.034   0.303
## season8      4.315e-02  3.010e-02   1.433   0.154
## season9      2.931e-02  3.010e-02   0.974   0.332
## season10     1.768e-02  3.010e-02   0.587   0.558
## season11     1.906e-02  3.009e-02   0.633   0.528
## season12     2.838e-02  3.009e-02   0.943   0.347
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07057 on 119 degrees of freedom
## Multiple R-squared:  0.03497,    Adjusted R-squared:  -0.06235
## F-statistic: 0.3593 on 12 and 119 DF,  p-value: 0.9748
```

```
plot(ts_stars)
lines(stars_tslm$fitted.values, col = "red")
```



```
detrend_stars = resid(stars_tslm)
plot(detrend_stars)
```



```
detrend_stars_lm = lm(detrend_stars ~ stars_recession_dummy)
summary(detrend_stars_lm)

##
## Call:
## lm(formula = detrend_stars ~ stars_recession_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.11605 -0.04819 -0.01401  0.05040  0.21376
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.007226   0.006128   1.179  0.24049
## stars_recession_dummy -0.050203   0.016153  -3.108  0.00231 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06514 on 130 degrees of freedom
## Multiple R-squared:  0.06917,    Adjusted R-squared:  0.06201
## F-statistic:  9.66 on 1 and 130 DF,  p-value: 0.002314
```

the growth rates plot shows that even during recession we don't see a dip in either. if there was constant account creation but a dip in reviews, that would mean the recession has an effect on num of reviews

## descriptive stats of reviews and stars

```
sd(df_rev_count$coredata.df_rev_m.)
```

```
## [1] 14205.68
```

```
mean(df_rev_count$coredata.df_rev_m.)
```

```
## [1] 15550.11
```

```
sd(log_rev_count)
```

```
## [1] 0.1694621
```

```
mean(log_rev_count)
```

```
## [1] 0.0357887
```

```
sd(df_stars$avg)
```

```
## [1] 0.06847152
```

```
mean(df_stars$avg)
```

```
## [1] 3.705981
```

## fred data

```
# real gdp
```

```
getSymbols("GDPC96", src = "FRED")
```

```
## As of 0.4-0, 'getSymbols' uses env=parent.frame() and
```

```
## auto.assign=TRUE by default.
```

```
##
```

```
## This behavior will be phased out in 0.5-0 when the call will
```

```
## default to use auto.assign=FALSE. getOption("getSymbols.env") and
```

```
## getOptions("getSymbols.auto.assign") are now checked for alternate defaults
```

```
##
```

```
## This message is shown once per session and may be disabled by setting
```


```
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for more details.
```

```
## [1] "GDPC96"
```


```
gdp = GDPC96
```

```
plot(gdp)
```


## gdp




```
gdp_growth = na.omit(diff(log(gdp)))  
plot(gdp_growth)
```



```
gdp_growth_subset = with(gdp_growth, gdp_growth[index(gdp_growth) >=
  "2006-04-01" & index(gdp_growth) < "2016-12-30", ])
gdp_growth_subset = ts(gdp_growth_subset, start = c(2006, 2),
  frequency = 4)
plot(gdp_growth_subset)
```




```
# create quarterly review growth rate
reviews_by_quarter = xts(review_counts_by_date$count(date),
  as.Date(review_counts_by_date$date, "%Y-%m-%d"))
df_rev_m_quarter = apply.quarterly(reviews_by_quarter, sum)
df_rev_quarter = data.frame(date = index(df_rev_m_quarter), coredata(df_rev_m_quarter))
# df_rev_quarter
plot(df_rev_quarter$date, df_rev_quarter$coredata.df_rev_m_quarter.,
  type = "l")
```



```
log_rev_quarter = diff(log(df_rev_quarter$coredata.df_rev_m_quarter.))  
log_rev_quarter = na.omit(log_rev_quarter)  
log_rev_quarter = ts(log_rev_quarter, start = c(2006, 2), freq = 4)  
plot(log_rev_quarter, type = "l")
```






```
# length(log_rev_quarter) length(gdp_growth_subset)

# var of gdp and user reviews growth rates
gdp_combined = cbind(log_rev_quarter, gdp_growth_subset)
select = VARselect(gdp_combined, lag.max = 12, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)
vm_gdp = VAR(gdp_combined, p = 12)
plot(vm_gdp$y)
```

## vm\_gdp\$y



```
summary(vm_gdp)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: log_rev_quarter, gdp_growth_subset
## Deterministic variables: const
## Sample size: 31
## Log Likelihood: 238.784
## Roots of the characteristic polynomial:
## 0.9987 0.9987 0.9921 0.9921 0.9906 0.9795 0.9795 0.9699 0.9612 0.9612 0.9576 0.9576 0.9327 0.9327 0.9327 0.9327
## Call:
## VAR(y = gdp_combined, p = 12)
##
##
## Estimation results for equation log_rev_quarter:
## =====
## log_rev_quarter = log_rev_quarter.l1 + gdp_growth_subset.l1 + log_rev_quarter.l2 + gdp_growth_subset.l2 +
##
##
```

	Estimate	Std. Error	t value	Pr(> t )
log_rev_quarter.l1	0.41495	0.36251	1.145	0.2960
gdp_growth_subset.l1	-0.69319	3.48625	-0.199	0.8490
log_rev_quarter.l2	0.28843	0.37771	0.764	0.4740
gdp_growth_subset.l2	4.10864	3.15933	1.300	0.2412
log_rev_quarter.l3	0.27193	0.32528	0.836	0.4352
gdp_growth_subset.l3	1.55990	3.26731	0.477	0.6500

```

## log_rev_quarter.l4      0.63886      0.20809      3.070      0.0219 *
## gdp_growth_subset.l4    1.46361      2.96026      0.494      0.6386
## log_rev_quarter.l5     -0.21532      0.23826     -0.904      0.4010
## gdp_growth_subset.l5   -0.43820      3.45264     -0.127      0.9032
## log_rev_quarter.l6     -0.16586      0.22996     -0.721      0.4979
## gdp_growth_subset.l6    4.69612      3.45845      1.358      0.2233
## log_rev_quarter.l7     -0.13165      0.24321     -0.541      0.6078
## gdp_growth_subset.l7   -4.04228      2.95912     -1.366      0.2209
## log_rev_quarter.l8      0.05838      0.24925      0.234      0.8226
## gdp_growth_subset.l8    5.04951      3.49119      1.446      0.1982
## log_rev_quarter.l9      0.11847      0.20277      0.584      0.5803
## gdp_growth_subset.l9    0.05217      3.36605      0.016      0.9881
## log_rev_quarter.l10     0.11208      0.13776      0.814      0.4470
## gdp_growth_subset.l10   6.34267      3.22213      1.968      0.0966 .
## log_rev_quarter.l11    -0.08093      0.15120     -0.535      0.6117
## gdp_growth_subset.l11  -1.75022      2.11397     -0.828      0.4394
## log_rev_quarter.l12     0.43781      0.17667      2.478      0.0479 *
## gdp_growth_subset.l12   2.53313      2.12131      1.194      0.2775
## const                   -0.16460      0.20382     -0.808      0.4502
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.04107 on 6 degrees of freedom
## Multiple R-Squared: 0.9823, Adjusted R-squared: 0.9117
## F-statistic: 13.9 on 24 and 6 DF, p-value: 0.001734
##
##
## Estimation results for equation gdp_growth_subset:
## =====
## gdp_growth_subset = log_rev_quarter.l1 + gdp_growth_subset.l1 + log_rev_quarter.l2 + gdp_growth_subset.l2
##
##
##               Estimate Std. Error t value Pr(>|t|)
## log_rev_quarter.l1 -0.036465   0.038485  -0.947   0.3800
## gdp_growth_subset.l1 -0.473658   0.370109  -1.280   0.2479
## log_rev_quarter.l2 -0.042996   0.040098  -1.072   0.3248
## gdp_growth_subset.l2 -0.403692   0.335402  -1.204   0.2741
## log_rev_quarter.l3 -0.005267   0.034533  -0.153   0.8838
## gdp_growth_subset.l3 -0.410470   0.346866  -1.183   0.2814
## log_rev_quarter.l4 -0.034836   0.022091  -1.577   0.1659
## gdp_growth_subset.l4 -0.611076   0.314269  -1.944   0.0998 .
## log_rev_quarter.l5 -0.004935   0.025294  -0.195   0.8518
## gdp_growth_subset.l5 -0.259610   0.366540  -0.708   0.5053
## log_rev_quarter.l6  0.026056   0.024413   1.067   0.3269
## gdp_growth_subset.l6 -0.470368   0.367157  -1.281   0.2474
## log_rev_quarter.l7 -0.014837   0.025820  -0.575   0.5864
## gdp_growth_subset.l7 -0.179878   0.314147  -0.573   0.5877
## log_rev_quarter.l8 -0.028504   0.026461  -1.077   0.3228
## gdp_growth_subset.l8 -0.567262   0.370633  -1.531   0.1768
## log_rev_quarter.l9 -0.002516   0.021527  -0.117   0.9108
## gdp_growth_subset.l9 -0.410645   0.357347  -1.149   0.2942
## log_rev_quarter.l10 -0.021819   0.014625  -1.492   0.1863
## gdp_growth_subset.l10 -0.189002   0.342069  -0.553   0.6006
## log_rev_quarter.l11 -0.034597   0.016052  -2.155   0.0746 .

```

```
## gdp_growth_subset.l11 0.090608 0.224424 0.404 0.7004
## log_rev_quarter.l12 -0.014337 0.018756 -0.764 0.4736
## gdp_growth_subset.l12 -0.408396 0.225203 -1.813 0.1197
## const 0.043626 0.021638 2.016 0.0904 .
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.00436 on 6 degrees of freedom
## Multiple R-Squared: 0.7602, Adjusted R-squared: -0.1992
## F-statistic: 0.7924 on 24 and 6 DF, p-value: 0.6886
##
##
## Covariance matrix of residuals:
##      log_rev_quarter gdp_growth_subset
## log_rev_quarter      0.0016865      1.157e-04
## gdp_growth_subset      0.0001157      1.901e-05
##
## Correlation matrix of residuals:
##      log_rev_quarter gdp_growth_subset
## log_rev_quarter      1.0000      0.6464
## gdp_growth_subset      0.6464      1.0000
# do a var on quarterly gdp growth and quarterly avg review
# score

recession_dummy_reviews_q = add_recession_dummy(df_rev_quarter$date)

reg_reviews = lm(df_rev_quarter$coredata.df_rev_m_quarter. ~
  recession_dummy_reviews_q)
summary(reg_reviews)

##
## Call:
## lm(formula = df_rev_quarter$coredata.df_rev_m_quarter. ~ recession_dummy_reviews_q)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53417 -28547  -3776   24826   78800
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      53903      6553   8.225 2.72e-10 ***
## recession_dummy_reviews_q -45589      16430  -2.775 0.00821 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39860 on 42 degrees of freedom
## Multiple R-squared: 0.1549, Adjusted R-squared: 0.1348
## F-statistic: 7.699 on 1 and 42 DF, p-value: 0.008211
# first reg is misleading
```

```
# look at growth rates
```

```
rec_dummy_rev_growth_q = recession_dummy_reviews_q[2:length(recession_dummy_reviews_q)]
```

```
reg_log_reviews = lm(log_rev_quarter ~ rec_dummy_rev_growth_q)
summary(reg_log_reviews)
```

```
##
## Call:
## lm(formula = log_rev_quarter ~ rec_dummy_rev_growth_q)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.44588 -0.13871 -0.02718  0.09270  0.83847
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.10999     0.03739   2.941  0.00535 **
## rec_dummy_rev_growth_q  0.05857     0.09268   0.632  0.53093
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2244 on 41 degrees of freedom
## Multiple R-squared:  0.009647,    Adjusted R-squared:  -0.01451
## F-statistic: 0.3994 on 1 and 41 DF,  p-value: 0.5309
```

```
# use lm
```

```
log_reviews_lm = lm(log_rev_quarter ~ rec_dummy_rev_growth_q)
summary(log_reviews_lm)
```

```
##
## Call:
## lm(formula = log_rev_quarter ~ rec_dummy_rev_growth_q)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.44588 -0.13871 -0.02718  0.09270  0.83847
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.10999     0.03739   2.941  0.00535 **
## rec_dummy_rev_growth_q  0.05857     0.09268   0.632  0.53093
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2244 on 41 degrees of freedom
## Multiple R-squared:  0.009647,    Adjusted R-squared:  -0.01451
## F-statistic: 0.3994 on 1 and 41 DF,  p-value: 0.5309
```

## examine social dynamics during recessions through dollar signs as well as review texts

do people eat at cheaper places? do people care more about overpriced food?

```
dollars_1_xts = xts(dollars_gbd_1$count(date)~, as.Date(dollars_gbd_1$date,
"%Y-%m-%d"))
df_d_1 = apply.monthly(dollars_1_xts, sum)
df_dollars_1 = data.frame(date = index(df_d_1), coredata(df_d_1))
# df_dollars_1


dollars_2_xts = xts(dollars_gbd_2$count(date)~, as.Date(dollars_gbd_2$date,
"%Y-%m-%d"))
df_d_2 = apply.monthly(dollars_2_xts, sum)
df_dollars_2 = data.frame(date = index(df_d_2), coredata(df_d_2))
# df_dollars_2

dollars_3_xts = xts(dollars_gbd_3$count(date)~, as.Date(dollars_gbd_3$date,
"%Y-%m-%d"))
df_d_3 = apply.monthly(dollars_3_xts, sum)
df_dollars_3 = data.frame(date = index(df_d_3), coredata(df_d_3))
# df_dollars_3


dollars_4_xts = xts(dollars_gbd_4$count(date)~, as.Date(dollars_gbd_4$date,
"%Y-%m-%d"))
df_d_4 = apply.monthly(dollars_4_xts, sum)
df_dollars_4 = data.frame(date = index(df_d_4), coredata(df_d_4))
# df_dollars_4

recession_dummy_dollars_m = add_recession_dummy(df_dollars_1$date)

plot(df_dollars_1, type = "l")
```




```
plot(df_dollars_2, type = "l")
```




```
plot(df_dollars_3, type = "l")
```





```
plot(df_dollars_4, type = "l")
```




```
df_dollars_1_dlog = as.data.frame(diff(log(df_dollars_1$coredata.df_d_1.)))
df_dollars_2_dlog = as.data.frame(diff(log(df_dollars_2$coredata.df_d_2.)))
df_dollars_3_dlog = as.data.frame(diff(log(df_dollars_3$coredata.df_d_3.)))
df_dollars_4_dlog = as.data.frame(diff(log(df_dollars_4$coredata.df_d_4.)))


# ts of dlogs
ts_dollar_1 = ts(df_dollars_1$coredata.df_d_1., start = c(2006,
2), freq = 12)
ts_dollar_2 = ts(df_dollars_2$coredata.df_d_2., start = c(2006,
2), freq = 12)
ts_dollar_3 = ts(df_dollars_3$coredata.df_d_3., start = c(2006,
2), freq = 12)
ts_dollar_4 = ts(df_dollars_4$coredata.df_d_4., start = c(2006,
2), freq = 12)

dollars_1_dlog = diff(log(ts_dollar_1))
dollars_2_dlog = diff(log(ts_dollar_2))
dollars_3_dlog = diff(log(ts_dollar_3))
dollars_4_dlog = diff(log(ts_dollar_4))


plot(dollars_1_dlog)
```




```
plot(dollars_2_dlog)
```



```
plot(dollars_3_dlog)
```



```
plot(dollars_4_dlog)
```




```
lm_dollars_recession_dummy = recession_dummy_dollars_m[2:length(recession_dummy_dollars_m)]
```

## seasonally adjust the data

```
# seasonally adjust data


# levels
tslm_d1 = tslm(ts_dollar_1 ~ trend + season)
# summary(tslm_d1)
tslm_d1_resid = resid(tslm_d1)
plot(tslm_d1_resid)
```



```
lm_d1_adj = lm(tslm_d1_resid ~ recession_dummy_dollars_m)
summary(lm_d1_adj)

##
## Call:
## lm(formula = tslm_d1_resid ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1932.8 -1098.0  -162.2   975.3  2613.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)         33.99    120.76   0.281   0.779
## recession_dummy_dollars_m -236.11    318.30  -0.742   0.460
##
## Residual standard error: 1284 on 130 degrees of freedom
## Multiple R-squared:  0.004215, Adjusted R-squared: -0.003445
## F-statistic: 0.5502 on 1 and 130 DF, p-value: 0.4596

tslm_d2 = tslm(ts_dollar_2 ~ trend + season)
# summary(tslm_d1)
tslm_d2_resid = resid(tslm_d2)
plot(tslm_d2_resid)
```




```
lm_d2_adj = lm(tslm_d2_resid ~ recession_dummy_dollars_m)
summary(lm_d2_adj)

##
## Call:
## lm(formula = tslm_d2_resid ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4262.1 -2293.1  -201.7   2170.2  5499.9
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)         71.32    259.47   0.275   0.784
## recession_dummy_dollars_m -495.52    683.92  -0.725   0.470
##
## Residual standard error: 2758 on 130 degrees of freedom
## Multiple R-squared:  0.004022, Adjusted R-squared:  -0.00364
## F-statistic: 0.5249 on 1 and 130 DF, p-value: 0.47

tslm_d3 = tslm(ts_dollar_3 ~ trend + season)
# summary(tslm_d1)
tslm_d3_resid = resid(tslm_d3)
plot(tslm_d3_resid)
```






```
lm_d3_adj = lm(tslm_d3_resid ~ recession_dummy_dollars_m)
summary(lm_d3_adj)
```

```
##
## Call:
## lm(formula = tslm_d3_resid ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -301.93  -84.56  -10.65   76.69  290.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)         14.40      10.99   1.310  0.19253
## recession_dummy_dollars_m -100.01      28.97  -3.453  0.00075 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 116.8 on 130 degrees of freedom
## Multiple R-squared:  0.084, Adjusted R-squared:  0.07695
## F-statistic: 11.92 on 1 and 130 DF, p-value: 0.0007496


tslm_d4 = tslm(ts_dollar_4 ~ trend + season)
# summary(tslm_d1)
tslm_d4_resid = resid(tslm_d4)
plot(tslm_d4_resid)
```



```
lm_d4_adj = lm(tslm_d4_resid ~ recession_dummy_dollars_m)
summary(lm_d4_adj)
```

```
##
## Call:
## lm(formula = tslm_d4_resid ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -158.13  -21.63    5.18   29.74  107.62
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)         4.168      4.353   0.958  0.3400
## recession_dummy_dollars_m -28.960     11.472  -2.524  0.0128 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 46.27 on 130 degrees of freedom
## Multiple R-squared:  0.04673,    Adjusted R-squared:  0.03939
## F-statistic: 6.372 on 1 and 130 DF,  p-value: 0.0128
```

```
# growth
tslm_d1_growth = tslm(dollars_1_dlog ~ trend + season)
# summary(tslm_d1)
tslm_d1_growth_resid = resid(tslm_d1_growth)
plot(tslm_d1_growth_resid)
```



```
lm_d1_growth_adj = lm(tslm_d1_growth_resid ~ lm_dollars_recession_dummy)
summary(lm_d1_growth_adj)
```

```
##
## Call:
## lm(formula = tslm_d1_growth_resid ~ lm_dollars_recession_dummy)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-1.15582	-0.06009	-0.00420	0.05984	0.88806

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.003396	0.018404	-0.185	0.854
lm_dollars_recession_dummy	0.023417	0.048325	0.485	0.629

```
##
## Residual standard error: 0.1948 on 129 degrees of freedom
## Multiple R-squared: 0.001817, Adjusted R-squared: -0.005921
## F-statistic: 0.2348 on 1 and 129 DF, p-value: 0.6288

tslm_d2_growth = tslm(dollars_2_dlog ~ trend + season)
# summary(tslm_d1)
tslm_d2_growth_resid = resid(tslm_d2_growth)
plot(tslm_d2_growth_resid)
```



```
lm_d2_growth_adj = lm(tslm_d2_growth_resid ~ lm_dollars_recession_dummy)
summary(lm_d2_growth_adj)
```

```
##
## Call:
## lm(formula = tslm_d2_growth_resid ~ lm_dollars_recession_dummy)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-0.71790	-0.04976	-0.00577	0.05448	0.55929

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.001668	0.013474	-0.124	0.902
lm_dollars_recession_dummy	0.011500	0.035381	0.325	0.746

```
##
## Residual standard error: 0.1426 on 129 degrees of freedom
## Multiple R-squared: 0.0008183, Adjusted R-squared: -0.006927
## F-statistic: 0.1056 on 1 and 129 DF, p-value: 0.7457
```


```
tslm_d3_growth = tslm(dollars_3_dlog ~ trend + season)
# summary(tslm_d1)
tslm_d3_growth_resid = resid(tslm_d3_growth)
plot(tslm_d3_growth_resid)
```



```
lm_d3_growth_adj = lm(tslm_d3_growth_resid ~ lm_dollars_recession_dummy)
summary(lm_d3_growth_adj)
```

```
##
## Call:
## lm(formula = tslm_d3_growth_resid ~ lm_dollars_recession_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.48967 -0.06243  0.00702  0.04723  0.72817
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -0.002172   0.013268  -0.164   0.870
## lm_dollars_recession_dummy  0.014978   0.034838   0.430   0.668
##
## Residual standard error: 0.1404 on 129 degrees of freedom
## Multiple R-squared:  0.001431, Adjusted R-squared:  -0.00631
## F-statistic: 0.1848 on 1 and 129 DF, p-value: 0.668

tslm_d4_growth = tslm(dollars_4_dlog ~ trend + season)
# summary(tslm_d1)
tslm_d4_growth_resid = resid(tslm_d4_growth)
plot(tslm_d4_growth_resid)
```



```
lm_d4_growth_adj = lm(tslm_d4_growth_resid ~ lm_dollars_recession_dummy)
summary(lm_d4_growth_adj)
```

```
##
## Call:
## lm(formula = tslm_d4_growth_resid ~ lm_dollars_recession_dummy)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.90447	-0.08604	-0.00312	0.09386	0.79209

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.001376	0.020201	0.068	0.946
lm_dollars_recession_dummy	-0.009485	0.053043	-0.179	0.858

```
##
## Residual standard error: 0.2138 on 129 degrees of freedom
## Multiple R-squared: 0.0002478, Adjusted R-squared: -0.007502
## F-statistic: 0.03198 on 1 and 129 DF, p-value: 0.8584
```


```
# percentages
```

```
ts_dollar_total = ts_dollar_1 + ts_dollar_2 + ts_dollar_3 + ts_dollar_4
ts_d1_per = ts_dollar_1/ts_dollar_total
ts_d2_per = ts_dollar_2/ts_dollar_total
ts_d3_per = ts_dollar_3/ts_dollar_total
ts_d4_per = ts_dollar_4/ts_dollar_total
```

```
plot(ts_d1_per)
```




```
plot(ts_d2_per)
```




```
plot(ts_d3_per)
```






```
plot(ts_d4_per)
```



```
tslm_d1_per = tslm(ts_d1_per ~ trend + season)
tslm_d1_per_resid = resid(tslm_d1_per)
plot(tslm_d1_per_resid)
```



```
lm_d1_per_adj = lm(tslm_d1_per_resid ~ recession_dummy_dollars_m)
summary(lm_d1_per_adj)
```

```
##
## Call:
## lm(formula = tslm_d1_per_resid ~ recession_dummy_dollars_m)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-0.087316	-0.005327	0.001185	0.009100	0.077087

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.0003379	0.0017737	0.191	0.849
recession_dummy_dollars_m	-0.0023478	0.0046751	-0.502	0.616

```
##
## Residual standard error: 0.01885 on 130 degrees of freedom
## Multiple R-squared: 0.001936, Adjusted R-squared: -0.005741
## F-statistic: 0.2522 on 1 and 130 DF, p-value: 0.6164
```

```
tslm_d2_per = tslm(ts_d2_per ~ trend + season)
tslm_d2_per_resid = resid(tslm_d2_per)
plot(tslm_d2_per_resid)
```



```
lm_d2_per_adj = lm(tslm_d2_per_resid ~ recession_dummy_dollars_m)
summary(lm_d2_per_adj)
```

```
##
## Call:
## lm(formula = tslm_d2_per_resid ~ recession_dummy_dollars_m)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-0.074900	-0.005344	0.000920	0.006557	0.043791

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.001535	0.001324	-1.159	0.24858
recession_dummy_dollars_m	0.010663	0.003491	3.055	0.00273 **

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01408 on 130 degrees of freedom
## Multiple R-squared:  0.06698,    Adjusted R-squared:  0.0598
## F-statistic: 9.332 on 1 and 130 DF,  p-value: 0.002732

tslm_d3_per = tslm(ts_d3_per ~ trend + season)
tslm_d3_per_resid = resid(tslm_d3_per)
plot(tslm_d3_per_resid)
```



```
lm_d3_per_adj = lm(tslm_d3_per_resid ~ recession_dummy_dollars_m)
summary(lm_d3_per_adj)
```

```
##
## Call:
## lm(formula = tslm_d3_per_resid ~ recession_dummy_dollars_m)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-0.033760	-0.006447	-0.001380	0.003443	0.098237

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.0007874	0.0013541	0.581	0.562
recession_dummy_dollars_m	-0.0054702	0.0035690	-1.533	0.128

```
##
## Residual standard error: 0.01439 on 130 degrees of freedom
## Multiple R-squared: 0.01775, Adjusted R-squared: 0.01019
## F-statistic: 2.349 on 1 and 130 DF, p-value: 0.1278
```

```
tslm_d4_per = tslm(ts_d4_per ~ trend + season)
tslm_d4_per_resid = resid(tslm_d4_per)
plot(tslm_d4_per_resid)
```



```
lm_d4_per_adj = lm(tslm_d4_per_resid ~ recession_dummy_dollars_m)
summary(lm_d4_per_adj)
```

```
##
## Call:
## lm(formula = tslm_d4_per_resid ~ recession_dummy_dollars_m)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.024128	-0.005074	-0.000451	0.003447	0.075637

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.0004095	0.0010212	0.401	0.689
recession_dummy_dollars_m	-0.0028453	0.0026917	-1.057	0.292

```
##
## Residual standard error: 0.01086 on 130 degrees of freedom
## Multiple R-squared: 0.008522, Adjusted R-squared: 0.0008948
## F-statistic: 1.117 on 1 and 130 DF, p-value: 0.2925
```

## introduce num of reviews into model

control for number of new reviews, to make sure that the recession dummy doesn't accidentally capture an effect from new review numbers, in other words if the number of reviews dropped it would make sense that the number of reviews for a certain amount of dollar signs drops too

```

# levels
lm_d1_adj_rev = lm(tslm_d1_resid ~ recession_dummy_dollars_m +
  df_rev_count$coredata.df_rev_m.)
summary(lm_d1_adj_rev)

##
## Call:
## lm(formula = tslm_d1_resid ~ recession_dummy_dollars_m + df_rev_count$coredata.df_rev_m.)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1829.6  -917.4  -307.4   1123.4   2837.2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -4.521e+02  1.852e+02  -2.441  0.016016 *
## recession_dummy_dollars_m      1.689e+02  3.291e+02   0.513  0.608659
## df_rev_count$coredata.df_rev_m.  2.751e-02  8.162e-03   3.370  0.000991 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1235 on 129 degrees of freedom
## Multiple R-squared:  0.0848, Adjusted R-squared:  0.07061
## F-statistic: 5.976 on 2 and 129 DF,  p-value: 0.003295

lm_d2_adj_rev = lm(tslm_d2_resid ~ recession_dummy_dollars_m +
  df_rev_count$coredata.df_rev_m.)
summary(lm_d2_adj_rev)

##
## Call:
## lm(formula = tslm_d2_resid ~ recession_dummy_dollars_m + df_rev_count$coredata.df_rev_m.)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4138.1 -1943.8  -620.9   2367.0   6147.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -977.68395  397.81376  -2.458  0.015313 *
## recession_dummy_dollars_m      378.55976  706.74989   0.536  0.593133
## df_rev_count$coredata.df_rev_m.    0.05937    0.01753   3.387  0.000938 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2653 on 129 degrees of freedom
## Multiple R-squared:  0.08534, Adjusted R-squared:  0.07116
## F-statistic: 6.018 on 2 and 129 DF,  p-value: 0.003171

lm_d3_adj_rev = lm(tslm_d3_resid ~ recession_dummy_dollars_m +
  df_rev_count$coredata.df_rev_m.)
summary(lm_d3_adj_rev)

##
## Call:

```


```
## lm(formula = tslm_d3_resid ~ recession_dummy_dollars_m + df_rev_count$coredata.df_rev_m.)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -315.824  -83.622   -8.919   77.962  280.690
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.403e+00  1.754e+01   0.194  0.84646
## recession_dummy_dollars_m    -9.085e+01  3.116e+01  -2.916  0.00418 **
## df_rev_count$coredata.df_rev_m.  6.221e-04  7.729e-04   0.805  0.42231
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 117 on 129 degrees of freedom
## Multiple R-squared:  0.08857,    Adjusted R-squared:  0.07444
## F-statistic: 6.268 on 2 and 129 DF,  p-value: 0.002523
lm_d4_adj_rev = lm(tslm_d4_resid ~ recession_dummy_dollars_m +
  df_rev_count$coredata.df_rev_m.)
summary(lm_d4_adj_rev)

##
## Call:
## lm(formula = tslm_d4_resid ~ recession_dummy_dollars_m + df_rev_count$coredata.df_rev_m.)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -145.400  -28.213    4.314   30.616  111.525
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      1.591e+01  6.836e+00   2.327  0.02152 *
## recession_dummy_dollars_m    -3.874e+01  1.214e+01  -3.190  0.00179 **
## df_rev_count$coredata.df_rev_m.  -6.644e-04  3.012e-04  -2.206  0.02919 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 45.6 on 129 degrees of freedom
## Multiple R-squared:  0.08137,    Adjusted R-squared:  0.06712
## F-statistic: 5.713 on 2 and 129 DF,  p-value: 0.004195
```

re examine stars, but subset by dollar signs

```
dollars_1_star_xts = xts(dollars_obd_1_star$stars, as.Date(dollars_obd_1_star$date,
  "%Y-%m-%d"))
df_d_1_star = apply.monthly(dollars_1_star_xts, sum)
df_dollars_1_star = data.frame(date = index(df_d_1_star), coredata(df_d_1_star))
df_dollars_1_star$avg = df_dollars_1_star$coredata.df_d_1_star./df_dollars_1$coredata.df_d_1.
plot(df_dollars_1_star$date, df_dollars_1_star$avg)
```






```
d1_star_lm = lm(df_dollars_1_star$avg ~ recession_dummy_dollars_m)
summary(d1_star_lm)
```

```
##
## Call:
## lm(formula = df_dollars_1_star$avg ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.26498 -0.03895 -0.01530  0.01052  0.53502
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.750694   0.009036  415.069  <2e-16 ***
## recession_dummy_dollars_m -0.044273   0.023818  -1.859   0.0653 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09606 on 130 degrees of freedom
## Multiple R-squared:  0.02589,    Adjusted R-squared:  0.0184
## F-statistic: 3.455 on 1 and 130 DF,  p-value: 0.06532

dollars_2_star_xts = xts(dollars_obd_2_star$stars, as.Date(dollars_obd_2_star$date,
"%Y-%m-%d"))
df_d_2_star = apply.monthly(dollars_2_star_xts, sum)
df_dollars_2_star = data.frame(date = index(df_d_2_star), coredata(df_d_2_star))
df_dollars_2_star$avg = df_dollars_2_star$coredata.df_d_2_star./df_dollars_2_star$coredata.df_d_2.
```

```
plot(df_dollars_2_star$date, df_dollars_2_star$avg)
```




```
d2_star_lm = lm(df_dollars_2_star$avg ~ recession_dummy_dollars_m)
summary(d2_star_lm)
```

```
##
## Call:
## lm(formula = df_dollars_2_star$avg ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.16317 -0.06408 -0.01288  0.06248  0.17477
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.671420   0.007485  490.516 < 2e-16 ***
## recession_dummy_dollars_m -0.076331   0.019728  -3.869 0.000172 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07956 on 130 degrees of freedom
## Multiple R-squared:  0.1033, Adjusted R-squared:  0.09637
## F-statistic: 14.97 on 1 and 130 DF, p-value: 0.000172

dollars_3_star_xts = xts(dollars_obd_3_star$stars, as.Date(dollars_obd_3_star$date,
"%Y-%m-%d"))
df_d_3_star = apply.monthly(dollars_3_star_xts, sum)
```


```
df_dollars_3_star = data.frame(date = index(df_d_3_star), coredata(df_d_3_star))
df_dollars_3_star$avg = df_dollars_3_star$coredata.df_d_3_star./df_dollars_3$coredata.df_d_3.
plot(df_dollars_3_star$date, df_dollars_3_star$avg)
```



```
d3_star_lm = lm(df_dollars_3_star$avg ~ recession_dummy_dollars_m)
summary(d3_star_lm)
```

```
##
## Call:
## lm(formula = df_dollars_3_star$avg ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.25080 -0.05478 -0.00522  0.03954  0.40509
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.800796   0.008157  465.945  <2e-16 ***
## recession_dummy_dollars_m -0.020898   0.021501  -0.972    0.333
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08671 on 130 degrees of freedom
## Multiple R-squared:  0.007215, Adjusted R-squared:  -0.000422
## F-statistic: 0.9447 on 1 and 130 DF, p-value: 0.3329
```

```
dollars_4_star_xts = xts(dollars_obd_4_star$stars, as.Date(dollars_obd_4_star$date,
"%Y-%m-%d"))
df_d_4_star = apply.monthly(dollars_4_star_xts, sum)
df_dollars_4_star = data.frame(date = index(df_d_4_star), coredata(df_d_4_star))
df_dollars_4_star$avg = df_dollars_4_star$coredata.df_d_4_star./df_dollars_4_star$coredata.df_d_4.
plot(df_dollars_4_star$date, df_dollars_4_star$avg)
```



```
d4_star_lm = lm(df_dollars_4_star$avg ~ recession_dummy_dollars_m)
summary(d4_star_lm)
```

```
##
## Call:
## lm(formula = df_dollars_4_star$avg ~ recession_dummy_dollars_m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.33494 -0.06462 -0.00311  0.06433  0.41506
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.98494    0.01040   383.33  <2e-16 ***
## recession_dummy_dollars_m  0.03506    0.02740    1.28   0.203
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1105 on 130 degrees of freedom
```

```
## Multiple R-squared:  0.01244,    Adjusted R-squared:  0.004841
## F-statistic: 1.637 on 1 and 130 DF,  p-value: 0.203
# looks like the stars dropping during a recession was only
# in 1 and 2 dollar signs restaurants
```

comment out for now, cuz takes too long to process

analyze word usages

```
buildCorpus = function(data, stem) {
  corpus = Corpus(VectorSource(data))
  corpus = tm_map(corpus, content_transformer(tolower))
  corpus = tm_map(corpus, PlainTextDocument)
  corpus = tm_map(corpus, removePunctuation)
  corpus = tm_map(corpus, removeWords, stopWords)
  if (stem == 1)
    corpus = tm_map(corpus, stemDocument)
  return(corpus)
}

buildWordCloud = function(corpus, pal, val, name) {
  wordcloud(corpus, max.words = 75, random.order = FALSE, colors = brewer.pal(val,
    pal), main = name)
}

stopWords = removePunctuation(stopwords("SMART"))

restaurant_reviews_rec = with(restaurant_reviews, restaurant_reviews[(restaurant_reviews$date >=
  "2007-12" & restaurant_reviews$date <= "2009-06"), ])

restaurant_reviews_norec = with(restaurant_reviews, restaurant_reviews[(restaurant_reviews$date >
  "2009-06" & restaurant_reviews$date <= "2011-12"), ])

# create corpuses
corpus_reviews_rec = buildCorpus(restaurant_reviews_rec$text,
  0)
dtm_rec = DocumentTermMatrix(corpus_reviews_rec)
tidy_rec = tidy(dtm_rec)
ap_sentiments <- tidy_rec %>% inner_join(get_sentiments("bing"),
  by = c(term = "word"))

ap_sentiments


## # A tibble: 436,171 × 4
##       document      term count sentiment
##       <chr>      <chr> <dbl>    <chr>
## 1 character(0)    died     1 negative
## 2 character(0) enthusiasm 1 positive
## 3 character(0)    fantastic 2 positive
## 4 character(0)     good     1 positive
## 5 character(0)    horrible 1 negative
## 6 character(0)     love     1 positive
```

```
## 7 character(0) recommendations 1 positive
## 8 character(0) good 1 positive
## 9 character(0) nice 1 positive
## 10 character(0) pure 1 positive
## # ... with 436,161 more rows
```

```
ap_sentiments %>% count(document, sentiment, wt = count) %>%
  ungroup() %>% spread(sentiment, n, fill = 0) %>% mutate(sentiment = positive -
  negative) %>% arrange(sentiment)
```

```
## # A tibble: 1 × 4
##       document negative positive sentiment
##       <chr>      <dbl>      <dbl>      <dbl>
## 1 character(0)  158714    346991    188277
```

```
ap_sentiments %>% count(sentiment, term, wt = count) %>% ungroup() %>%
  filter(n >= 2000) %>% mutate(n = ifelse(sentiment == "negative",
  -n, n)) %>% mutate(term = reorder(term, n)) %>% ggplot(aes(term,
  n, fill = sentiment)) + geom_bar(stat = "identity") + theme(axis.text.x = element_text(angle = 90,
  hjust = 1)) + ylab("Contribution to sentiment")
```



```
# negative: 158714 positive: 346991 percent negative: 31.4%
# cheap: 3rd highest negative word expensive: 7th
```

```
corpus_reviews_norec = buildCorpus(restaurant_reviews_norec$text,
  0)
```

```

dtm_norec = DocumentTermMatrix(corpus_reviews_norec)
tidy_norec = tidy(dtm_norec)
ap_sentiments <- tidy_norec %>% inner_join(get_sentiments("bing"),
  by = c(term = "word"))

ap_sentiments


## # A tibble: 2,121,102 × 4
##       document      term count sentiment
##       <chr>      <chr> <dbl>      <chr>
## 1 character(0) amazingly     1 positive
## 2 character(0) awesome      1 positive
## 3 character(0) fast         2 positive
## 4 character(0) fucking      1 negative
## 5 character(0) great        1 positive
## 6 character(0) holy         1 positive
## 7 character(0) nice         1 positive
## 8 character(0) shit         1 negative
## 9 character(0) weak         1 negative
## 10 character(0) work         1 positive
## # ... with 2,121,092 more rows

ap_sentiments %>% count(document, sentiment, wt = count) %>%
  ungroup() %>% spread(sentiment, n, fill = 0) %>% mutate(sentiment = positive -
  negative) %>% arrange(sentiment)

## # A tibble: 1 × 4
##       document negative positive sentiment
##       <chr>      <dbl>      <dbl>      <dbl>
## 1 character(0)  733329  1731938   998609

ap_sentiments %>% count(sentiment, term, wt = count) %>% ungroup() %>%
  filter(n >= 7500) %>% mutate(n = ifelse(sentiment == "negative",
  -n, n)) %>% mutate(term = reorder(term, n)) %>% ggplot(aes(term,
  n, fill = sentiment)) + geom_bar(stat = "identity") + theme(axis.text.x = element_text(angle = 90,
  hjust = 1)) + ylab("Contribution to sentiment")


```



# negative: 733329 positive: 1731938 percent negative: 29.75%  
 # cheap: 5th expensive: 8th

```
buildWordCloud(corpus_reviews_rec, "Spectral", 8, "test")
```








## connect with restaurants

```
# add x axis with dates
rest_exp_dates = seq(as.Date("2005/12/01"), by = "quarter", length.out = 44)

plot(rest_exp_dates, restaurant_expenditures$real_exp, type = "l")
```



```
test_stationary(restaurant_expenditures$real_exp)
```

```
## Warning in kpss.test(t): p-value smaller than printed p-value
```


```
##
## KPSS Test for Level Stationarity
##
## data: t
## KPSS Level = 1.7047, Truncation lag parameter = 1, p-value = 0.01
##
##
## Augmented Dickey-Fuller Test
##
## data: t
## Dickey-Fuller = -2.4338, Lag order = 3, p-value = 0.4021
## alternative hypothesis: stationary
```

```
rest_real_exp_diff_log = diff(log(restaurant_expenditures$real_exp))
```

```
test_stationary(rest_real_exp_diff_log)
```


```
##
## KPSS Test for Level Stationarity
##
## data: t
## KPSS Level = 0.69776, Truncation lag parameter = 1, p-value =
## 0.01375
##
```

```
##
## Augmented Dickey-Fuller Test
##
## data: t
## Dickey-Fuller = -2.2003, Lag order = 3, p-value = 0.4946
## alternative hypothesis: stationary
rest_exp_dates_diff = rest_exp_dates[2:length(rest_exp_dates)]
plot(rest_exp_dates_diff, rest_real_exp_diff_log, type = "l")
```



```
# create var + granger causality for real exp and gdp
gdp_exp_combined = cbind(rest_real_exp_diff_log, gdp_growth_subset)
select = VARselect(gdp_exp_combined, lag.max = 4, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)
vm_gdp_exp = VAR(gdp_exp_combined, select$select[1])
plot(vm_gdp_exp$y)
```

## vm\_gdp\_exp\$y



```
summary(vm_gdp_exp)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: rest_real_exp_diff_log, gdp_growth_subset
## Deterministic variables: const
## Sample size: 40
## Log Likelihood: 309.26
## Roots of the characteristic polynomial:
## 0.829 0.7485 0.7485 0.7006 0.5467 0.5467
## Call:
## VAR(y = gdp_exp_combined, p = select$select[1])
##
##
## Estimation results for equation rest_real_exp_diff_log:
## =====
## rest_real_exp_diff_log = rest_real_exp_diff_log.l1 + gdp_growth_subset.l1 + rest_real_exp_diff_log.l2
##
##               Estimate Std. Error t value Pr(>|t|)
## rest_real_exp_diff_log.l1 -0.3061642  0.2190439  -1.398  0.17152
## gdp_growth_subset.l1      0.7168516  0.2564126   2.796  0.00857 **
## rest_real_exp_diff_log.l2  0.1645218  0.1996566   0.824  0.41584
## gdp_growth_subset.l2      0.1951121  0.2662260   0.733  0.46880
## rest_real_exp_diff_log.l3  0.1923951  0.1868020   1.030  0.31053
## gdp_growth_subset.l3      0.1265239  0.2690216   0.470  0.64123
```

```

## const                -0.0001735  0.0013498  -0.129  0.89852
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.007011 on 33 degrees of freedom
## Multiple R-Squared:  0.4144, Adjusted R-squared:  0.3079
## F-statistic: 3.892 on 6 and 33 DF, p-value: 0.004777
##
##
## Estimation results for equation gdp_growth_subset:
## =====
## gdp_growth_subset = rest_real_exp_diff_log.l1 + gdp_growth_subset.l1 + rest_real_exp_diff_log.l2 + g
##
##               Estimate Std. Error t value Pr(>|t|)
## rest_real_exp_diff_log.l1 -0.095202  0.191103  -0.498  0.6217
## gdp_growth_subset.l1      0.469106  0.223705   2.097  0.0437 *
## rest_real_exp_diff_log.l2  0.070094  0.174189   0.402  0.6900
## gdp_growth_subset.l2     -0.008452  0.232267  -0.036  0.9712
## rest_real_exp_diff_log.l3  0.323437  0.162974   1.985  0.0556 .
## gdp_growth_subset.l3     -0.292756  0.234706  -1.247  0.2211
## const                    0.001641  0.001178   1.393  0.1729
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.006116 on 33 degrees of freedom
## Multiple R-Squared:  0.2667, Adjusted R-squared:  0.1333
## F-statistic:      2 on 6 and 33 DF, p-value: 0.09384
##
##
## Covariance matrix of residuals:
##               rest_real_exp_diff_log gdp_growth_subset
## rest_real_exp_diff_log      4.915e-05      2.948e-05
## gdp_growth_subset          2.948e-05      3.741e-05
##
## Correlation matrix of residuals:
##               rest_real_exp_diff_log gdp_growth_subset
## rest_real_exp_diff_log      1.0000      0.6874
## gdp_growth_subset          0.6874      1.0000
grangertest(rest_real_exp_diff_log ~ gdp_growth_subset[1:length(gdp_growth_subset)],
  order = select$select[1])

## Granger causality test
##
## Model 1: rest_real_exp_diff_log ~ Lags(rest_real_exp_diff_log, 1:3) + Lags(gdp_growth_subset[1:length
## Model 2: rest_real_exp_diff_log ~ Lags(rest_real_exp_diff_log, 1:3)
##   Res.Df Df      F Pr(>F)
## 1      33
## 2      36 -3 2.8912 0.05002 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

grangertest(gdp_growth_subset[1:length(gdp_growth_subset)] ~
  rest_real_exp_diff_log, order = select$select[1])

## Granger causality test
##
## Model 1: gdp_growth_subset[1:length(gdp_growth_subset)] ~ Lags(gdp_growth_subset[1:length(gdp_growth_subset)]
## Model 2: gdp_growth_subset[1:length(gdp_growth_subset)] ~ Lags(gdp_growth_subset[1:length(gdp_growth_subset)]
##   Res.Df Df       F Pr(>F)
## 1      33
## 2      36 -3 1.3863 0.2642

rec_exp_diff_log_dummy = add_recession_dummy(rest_exp_dates_diff)
lm_rest_real_exp_diff_log = lm(rest_real_exp_diff_log ~ rec_exp_diff_log_dummy)
summary(lm_rest_real_exp_diff_log)

##
## Call:
## lm(formula = rest_real_exp_diff_log ~ rec_exp_diff_log_dummy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0158435 -0.0038443 -0.0000908  0.0033176  0.0144583
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.005997   0.001061   5.654 1.34e-06 ***
## rec_exp_diff_log_dummy -0.015065   0.002629  -5.730 1.05e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.006364 on 41 degrees of freedom
## Multiple R-squared:  0.4447, Adjusted R-squared:  0.4312
## F-statistic: 32.84 on 1 and 41 DF, p-value: 1.046e-06

yes gdp granger causes restaurant expenditures and recession dummy


```

## connect num of reviews with restaurants

```

rev_exp_combined = cbind(rest_real_exp_diff_log, log_rev_quarter)
select = VARselect(rev_exp_combined, lag.max = 4, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)
vm_rev_exp = VAR(rev_exp_combined, select$select[1])
plot(vm_rev_exp$y)

```



```
summary(vm_rev_exp)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: rest_real_exp_diff_log, log_rev_quarter
## Deterministic variables: const
## Sample size: 39
## Log Likelihood: 190.353
## Roots of the characteristic polynomial:
## 0.9262 0.9126 0.868 0.868 0.5991 0.5991 0.5268 0.5268
## Call:
## VAR(y = rev_exp_combined, p = select$select[1])
##
##
## Estimation results for equation rest_real_exp_diff_log:
## =====
## rest_real_exp_diff_log = rest_real_exp_diff_log.l1 + log_rev_quarter.l1 + rest_real_exp_diff_log.l2 + log_rev_quarter.l2 + rest_real_exp_diff_log.l3 + log_rev_quarter.l3
##
##               Estimate Std. Error t value Pr(>|t|)
## rest_real_exp_diff_log.l1  0.2010747  0.1805061   1.114   0.2741
## log_rev_quarter.l1        -0.0077785  0.0088976  -0.874   0.3889
## rest_real_exp_diff_log.l2  0.2249583  0.1755076   1.282   0.2097
## log_rev_quarter.l2        -0.0041271  0.0076756  -0.538   0.5948
## rest_real_exp_diff_log.l3  0.1365037  0.1726607   0.791   0.4354
## log_rev_quarter.l3         0.0005648  0.0077776   0.073   0.9426
```



```

## rest_real_exp_diff_log.l4 -0.0140669  0.1650931  -0.085   0.9327
## log_rev_quarter.l4      -0.0107478  0.0078840  -1.363   0.1830
## const                   0.0044418  0.0025252   1.759   0.0888 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.007607 on 30 degrees of freedom
## Multiple R-Squared:  0.3578, Adjusted R-squared:  0.1866
## F-statistic:  2.09 on 8 and 30 DF,  p-value: 0.06888
##
##
## Estimation results for equation log_rev_quarter:
## =====
## log_rev_quarter = rest_real_exp_diff_log.l1 + log_rev_quarter.l1 + rest_real_exp_diff_log.l2 + log_r
##
##               Estimate Std. Error t value Pr(>|t|)
## rest_real_exp_diff_log.l1 -0.87806    1.83748  -0.478  0.63621
## log_rev_quarter.l1       0.12723    0.09057   1.405  0.17040
## rest_real_exp_diff_log.l2  3.28692    1.78660   1.840  0.07572 .
## log_rev_quarter.l2       0.07675    0.07813   0.982  0.33380
## rest_real_exp_diff_log.l3  0.73014    1.75762   0.415  0.68080
## log_rev_quarter.l3      -0.16649    0.07917  -2.103  0.04397 *
## rest_real_exp_diff_log.l4 -5.18564    1.68058  -3.086  0.00434 **
## log_rev_quarter.l4       0.54755    0.08026   6.823 1.44e-07 ***
## const                   0.03019    0.02571   1.174  0.24945
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.07744 on 30 degrees of freedom
## Multiple R-Squared:  0.8139, Adjusted R-squared:  0.7642
## F-statistic:  16.4 on 8 and 30 DF,  p-value: 5.127e-09
##
##
## Covariance matrix of residuals:
##               rest_real_exp_diff_log log_rev_quarter
## rest_real_exp_diff_log      5.787e-05      0.0001152
## log_rev_quarter            1.152e-04      0.0059970
##
## Correlation matrix of residuals:
##               rest_real_exp_diff_log log_rev_quarter
## rest_real_exp_diff_log      1.0000      0.1955
## log_rev_quarter            0.1955      1.0000
grangertest(rest_real_exp_diff_log ~ log_rev_quarter, order = select$select[1])

## Granger causality test
##
## Model 1: rest_real_exp_diff_log ~ Lags(rest_real_exp_diff_log, 1:4) + Lags(log_rev_quarter, 1:4)
## Model 2: rest_real_exp_diff_log ~ Lags(rest_real_exp_diff_log, 1:4)
##   Res.Df Df       F Pr(>F)
## 1      30
## 2      34 -4 0.9082 0.4718

```

```

grangertest(log_rev_quarter ~ rest_real_exp_diff_log, order = select$select[1])


## Granger causality test
##
## Model 1: log_rev_quarter ~ Lags(log_rev_quarter, 1:4) + Lags(rest_real_exp_diff_log, 1:4)
## Model 2: log_rev_quarter ~ Lags(log_rev_quarter, 1:4)
##   Res.Df Df       F    Pr(>F)
## 1      30
## 2      34 -4 2.8674 0.04005 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# looks like log_rev_quarter is granger caused by
# rest_real_exp_diff_log
log_rev_exp_lm = lm(log_rev_quarter ~ rec_dummy_rev_growth_q +
  rest_real_exp_diff_log)
summary(log_rev_exp_lm)


##
## Call:
## lm(formula = log_rev_quarter ~ rec_dummy_rev_growth_q + rest_real_exp_diff_log)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.48972 -0.13521 -0.02083  0.10232  0.79866
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.131999   0.047314   2.790  0.00804 **
## rec_dummy_rev_growth_q 0.007791   0.114349   0.068  0.94602
## rest_real_exp_diff_log -3.875810   5.061899  -0.766  0.44836
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2255 on 40 degrees of freedom
## Multiple R-squared:  0.02395,    Adjusted R-squared:  -0.02485
## F-statistic: 0.4908 on 2 and 40 DF,  p-value: 0.6158

plot(log_rev_exp_lm$fitted, type = "l")

```



```
plot(log_rev_quarter)
```




```
# gdp IV

# two stage regression
stage1 = lm(rest_real_exp_diff_log ~ gdp_growth_subset)
summary(stage1)

##
## Call:
## lm(formula = rest_real_exp_diff_log ~ gdp_growth_subset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0120421 -0.0052213  0.0003529  0.0046706  0.0125832
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.0004049  0.0010428   0.388    0.7
## gdp_growth_subset 0.9357058  0.1460001   6.409 1.13e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.006036 on 41 degrees of freedom
## Multiple R-squared:  0.5005, Adjusted R-squared:  0.4883
## F-statistic: 41.07 on 1 and 41 DF, p-value: 1.134e-07

stage1_fitted = fitted(stage1)
```

```
# var for possible iv
iv_combined = cbind(stage1_fitted, log_rev_quarter)
select = VARselect(iv_combined, lag.max = 4, type = c("const",
  "trend", "both", "none"), season = NULL, exogen = NULL)
vm_iv = VAR(iv_combined, select$select[1])
plot(vm_iv$y)
```



```
summary(vm_iv)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: stage1_fitted, log_rev_quarter
## Deterministic variables: const
## Sample size: 39
## Log Likelihood: 202.729
## Roots of the characteristic polynomial:
## 0.91 0.8766 0.8609 0.8609 0.7678 0.7678 0.6853 0.6853
## Call:
## VAR(y = iv_combined, p = select$select[1])
##
##
## Estimation results for equation stage1_fitted:
## =====
## stage1_fitted = stage1_fitted.l1 + log_rev_quarter.l1 + stage1_fitted.l2 + log_rev_quarter.l2 + stag
##
```

```

##               Estimate Std. Error t value Pr(>|t|)
## stage1_fitted.l1    0.487351   0.175025   2.784   0.0092 **
## log_rev_quarter.l1  0.007386   0.006431   1.148   0.2599
## stage1_fitted.l2   -0.140392   0.197013  -0.713   0.4816
## log_rev_quarter.l2 -0.005684   0.005587  -1.017   0.3171
## stage1_fitted.l3    0.203458   0.193332   1.052   0.3010
## log_rev_quarter.l3 -0.005723   0.005737  -0.998   0.3265
## stage1_fitted.l4   -0.226631   0.167025  -1.357   0.1849
## log_rev_quarter.l4 -0.007732   0.006001  -1.288   0.2074
## const              0.004077   0.001809   2.253   0.0317 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.005443 on 30 degrees of freedom
## Multiple R-Squared:  0.3944, Adjusted R-squared:  0.2329
## F-statistic: 2.442 on 8 and 30 DF, p-value: 0.03638
##
##
## Estimation results for equation log_rev_quarter:
## =====
## log_rev_quarter = stage1_fitted.l1 + log_rev_quarter.l1 + stage1_fitted.l2 + log_rev_quarter.l2 + stage1_fitted.l3 + log_rev_quarter.l3 + stage1_fitted.l4 + log_rev_quarter.l4 + const
##
##               Estimate Std. Error t value Pr(>|t|)
## stage1_fitted.l1   -3.07255    2.48728  -1.235   0.22630
## log_rev_quarter.l1  0.23450    0.09139   2.566   0.01553 *
## stage1_fitted.l2    4.00838    2.79976   1.432   0.16257
## log_rev_quarter.l2  0.11229    0.07940   1.414   0.16757
## stage1_fitted.l3    3.81243    2.74744   1.388   0.17547
## log_rev_quarter.l3 -0.24877    0.08153  -3.051   0.00474 **
## stage1_fitted.l4   -6.21185    2.37360  -2.617   0.01376 *
## log_rev_quarter.l4  0.51038    0.08528   5.985 1.46e-06 ***
## const              0.02798    0.02571   1.088   0.28519
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 0.07735 on 30 degrees of freedom
## Multiple R-Squared:  0.8143, Adjusted R-squared:  0.7648
## F-statistic: 16.44 on 8 and 30 DF, p-value: 4.95e-09
##
##
## Covariance matrix of residuals:
##               stage1_fitted log_rev_quarter
## stage1_fitted    2.962e-05   -1.602e-05
## log_rev_quarter  -1.602e-05    5.982e-03
##
## Correlation matrix of residuals:
##               stage1_fitted log_rev_quarter
## stage1_fitted    1.00000   -0.03807
## log_rev_quarter  -0.03807    1.00000

```

```
grangertest(stage1_fitted ~ log_rev_quarter, order = select$select[1])
```

```
## Granger causality test
```

```
##
```

```
## Model 1: stage1_fitted ~ Lags(stage1_fitted, 1:4) + Lags(log_rev_quarter, 1:4)
```

```
## Model 2: stage1_fitted ~ Lags(stage1_fitted, 1:4)
```

```
##   Res.Df Df       F Pr(>F)
```

```
## 1      30
```

```
## 2      34 -4 2.5951 0.0563 .
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
grangertest(log_rev_quarter ~ stage1_fitted, order = select$select[1])
```

```
## Granger causality test
```

```
##
```

```
## Model 1: log_rev_quarter ~ Lags(log_rev_quarter, 1:4) + Lags(stage1_fitted, 1:4)
```

```
## Model 2: log_rev_quarter ~ Lags(log_rev_quarter, 1:4)
```

```
##   Res.Df Df       F Pr(>F)
```

```
## 1      30
```

```
## 2      34 -4 2.8927 0.03881 *
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
old stuff
```

```
sources: fred yelp https://www.bea.gov/iTable/iTable.cfm?reqid=9&step=1&acrdrn=2#reqid=9&step=1&isuri=1&904=2004&903=64&906=q&905=2016&910=x&911=0
```